Agentic AI Interview – Experienced Hire

INTERVIEW SUCCESS PATTERNS
- Ask clarifying questions first
- Explain trade-offs explicitly
- Think out loud
- Discuss failure modes proactively

INTERVIEW CONTEXT
- **Target Roles**
  - Backend Engineers
  - ML Engineers
  - Platform Engineers
  - AI Engineers (3+ years)
- **Interview Focus**
  - System design thinking
  - Production experience
  - Failure handling and trade-offs

Q1: Multi Agent vs Single Agent
- **Interview Question**
  - When would you choose a single agent instead of a multi agent system?
- **What the Interviewer Is Testing**
  - Context driven architecture decisions
  - Trade-off thinking
  - Production cost and latency awareness
- **Strong Answer Signals**
  - Clarifies constraints before answering
  - Mentions latency and orchestration cost
  - Explains failure domain isolation
  - Gives scenario-based justification
- **Weak Answer Signals**
  - Multi agent is always better
  - Single agent is simpler so start there
  - No cost or latency discussion
- **Follow Up Questions**
  - How would you observe failures here?
  - What changes at 10x traffic?
  - How would you debug this in prod?

Q2: Planner vs Executor Roles
- **Interview Question**
  - Why separate planner and executor agents?
- **What the Interviewer Is Testing**
  - Failure isolation understanding
  - Control and observability mindset
- **Strong Answer Signals**
  - Planner owns state and decisions
  - Executors are constrained and replaceable
  - Mentions retry and substitution logic
- **Weak Answer Signals**
  - Just a design pattern
  - Cleaner code explanation only
- **Follow Up Questions**
  - What happens if the planner crashes?
  - How do you replay execution?

Q3: Memory Design Decisions
- **Interview Question**
  - How do you design memory for an agentic system?
- **What the Interviewer Is Testing**
  - System design thinking
  - Privacy and compliance awareness
  - Failure mode handling
- **Strong Answer Signals**
  - Separates short term and long term memory
  - Explains when NOT to store memory
  - Mentions audit and replay needs
- **Weak Answer Signals**
  - Just naming a vector database
  - No privacy or retention discussion
- **Follow Up Questions**
  - What if the memory store is down?
  - How do you handle sensitive data?

Q4: Failure Handling and Retries
- **Interview Question**
  - An external API call fails. What happens next?
- **What the Interviewer Is Testing**
  - Production readiness
  - Idempotency awareness
  - Operational experience
- **Strong Answer Signals**
  - Partial failure handling
  - Safe retries via idempotency
  - Timeouts and backoff strategy
  - Observability and tracing
- **Weak Answer Signals**
  - Retry three times with try catch
  - No observability discussion
- **Follow Up Questions**
  - How do you prevent cascading failures?
  - How would you debug this at 3 AM?

Q5: Evaluation and Human in the Loop
- **Interview Question**
  - How do you know your agent is working correctly?
- **What the Interviewer Is Testing**
  - Evaluation as a continuous process
  - Risk based automation decisions
- **Strong Answer Signals**
  - Automated evaluation plus human approval gates
  - Mentions regression and monitoring
  - Explains which actions require humans
- **Weak Answer Signals**
  - Test with a few examples
  - Human reviews everything
- **Follow Up Questions**
  - How do you detect drift?
  - How do you measure quality over time?

Join Bootcamp for Structured Learning:
https://bootcamp.nachiketh.in